

Book

Recommendation

System

Using Goodreads

Interactions Dataset

Berkehan Göktürk
Azra İrem Dalhasanoğlu
Beyzanur Elma



Dataset

In this study, we utilized the **Goodreads Interactions Dataset to develop a book recommendation system**. This dataset captures the interactions of registered Goodreads users with books – such as reading, rating, and shelving activities – and contains both user-item historical interaction data as well as metadata related to the books.

Dataset Analysis, Cleaning, and Merging

The dataset consists of two main files: interactions and books metadata. It contains numerous features such as user_id, book_id, publication_year, num_pages, and many others. In line with our project objectives, we **preprocessed the data** by removing irrelevant features, handling missing values either by imputation or deletion, and correcting characters to comply with UTF-8 encoding standards. Finally, we merged the two files to create a unified dataset, which served as the input for model training.

Dataset Analysis, Cleaning, and Merging

```
import pandas as pd
import re
# setting and constants
INPUT_BOOKS = "data/raw/books_metadata_large.csv"
INPUT_INTERACTIONS = "data/raw/interactions_large.csv"
OUTPUT_FILE = "data/processed/cleaned.csv"

# columns to keep
COLS_BOOK = [
    "book_id", "title", "average_rating", "ratings_count",
    "publication_year", "num_pages"
]
COLS_INTERACTION = ["user_id", "book_id", "rating"]

# regex for titles containing only Latin characters
latin_regex = re.compile(r'^[A-Za-z0-9\s\.,:;!?"-]+$')

# load the data
print("Loading data...")
# Load books metadata
df_books = pd.read_csv(INPUT_BOOKS, usecols=COLS_BOOK)
# Load user interactions
df_interactions = pd.read_csv(INPUT_INTERACTIONS, usecols=COLS_INTERACTION)
```

load the data

Dataset Analysis, Cleaning, and Merging

```
# filtering
print("Filtering books...")
# filtering by Latin characters
df_books = df_books[df_books['title'].apply(lambda x: bool(latin_regex.match(str(x))))]
# keep Latin titles
# year filter
df_books = df_books[(df_books['publication_year'] >= 1500) & (df_books['publication_year'] <= 2025)]
print("Filtering interactions...")
# rating filter
df_interactions = df_interactions[(df_interactions['rating'] >= 1) & (df_interactions['rating'] <= 5)]

# merging
print("Merging tables...")
# merge datasets
df_merged = pd.merge(df_interactions, df_books, on="book_id", how="inner")

# cleaning
print("Performing final cleaning...")
# remove missing data
df_merged = df_merged.dropna()
# remove duplicates
df_merged = df_merged.drop_duplicates(subset=['user_id', 'book_id'])
# remove invalid values
df_merged = df_merged[(df_merged['num_pages'] > 0) & (df_merged['ratings_count'] >= 0)]
```

filtering by Latin characters

filtering by year

filtering by rating

merge the tables

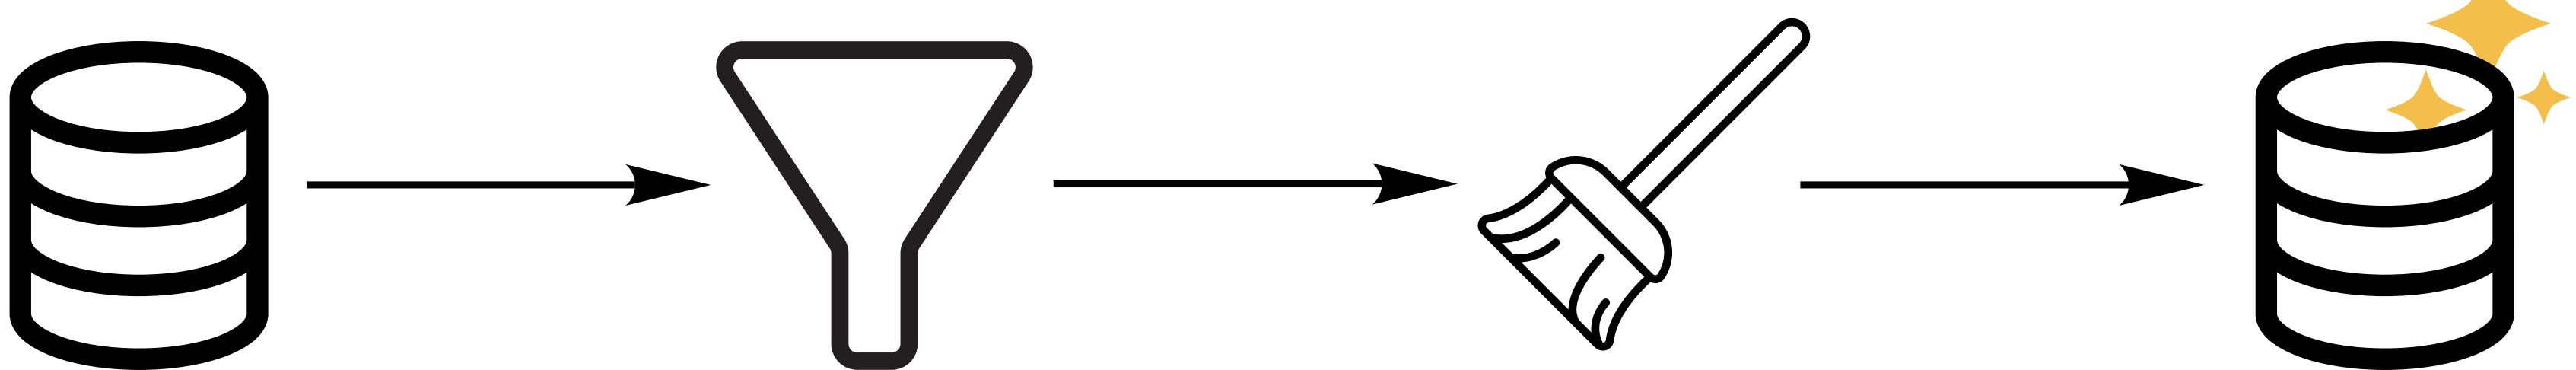
cleaning

Merge* interactions with books on 'book_id'.
Cleaning* remove missing data, duplications and outliers

the amount of raw data

interactions_large.csv : 2734350

books_metadata_large.csv : 36514



the amount of cleaned data

901670

Model Development: Approaches, Experiments, and Results

Model Development

- Read the dataset
- Classification
- Feature selection
- Split into training and test sets
- **Model creating and training**
- Predictions and evaluation

Model creating and training

- Logistic Regression
- Random Forest
- Multilayer Perceptron
- Boosting
- Stacking Ensemble
- SVD

try 1: Logistic Regression

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
import matplotlib.pyplot as plt
from sklearn.metrics import RocCurveDisplay
from sklearn.metrics import (accuracy_score, classification_report,precision_score,recall_score,f1_score)
from sklearn.metrics import roc_auc_score

# Read the dataset (relative path)
dataset = pd.read_csv("data/processed/cleaned.csv") read the dataset

# Binary classification: books with rating > 3 are considered popular
popularityBound = 3
dataset['label'] = (dataset['rating'] > popularityBound).astype(int) classification

# Select features
features = ['publication_year','num_pages','ratings_count']
print(dataset[['title','rating','label']].head()) # optional feature selection

D = dataset[features] # Feature matrix
y = dataset['label'] # Target variable

# Split the dataset into 70% train and 30% test
D_train, D_test, y_train, y_test = train_test_split(D, y, test_size=0.3, random_state=42) dataset splitting
```

try 1: Logistic Regression

```
# Scaler for feature normalization
scaler = StandardScaler()
D_train = scaler.fit_transform(D_train)
D_test = scaler.transform(D_test)

# Logistic Regression MODEL
model = LogisticRegression()
model.fit(D_train, y_train)

# Prediction and evaluation
y_pred = model.predict(D_test)

acc = accuracy_score(y_test, y_pred)
auc = roc_auc_score(y_test, model.predict_proba(D_test)[:,1])
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall   : {rec:.4f}")
print(f"F1-score : {f1:.4f}")
print(f"ROC-AUC  : {auc:.4f}")

print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

normalization

logistic regression model

prediction and accuracy

Normalization* to ensure all features contribute equally to the model.

try 1: Logistic Regression

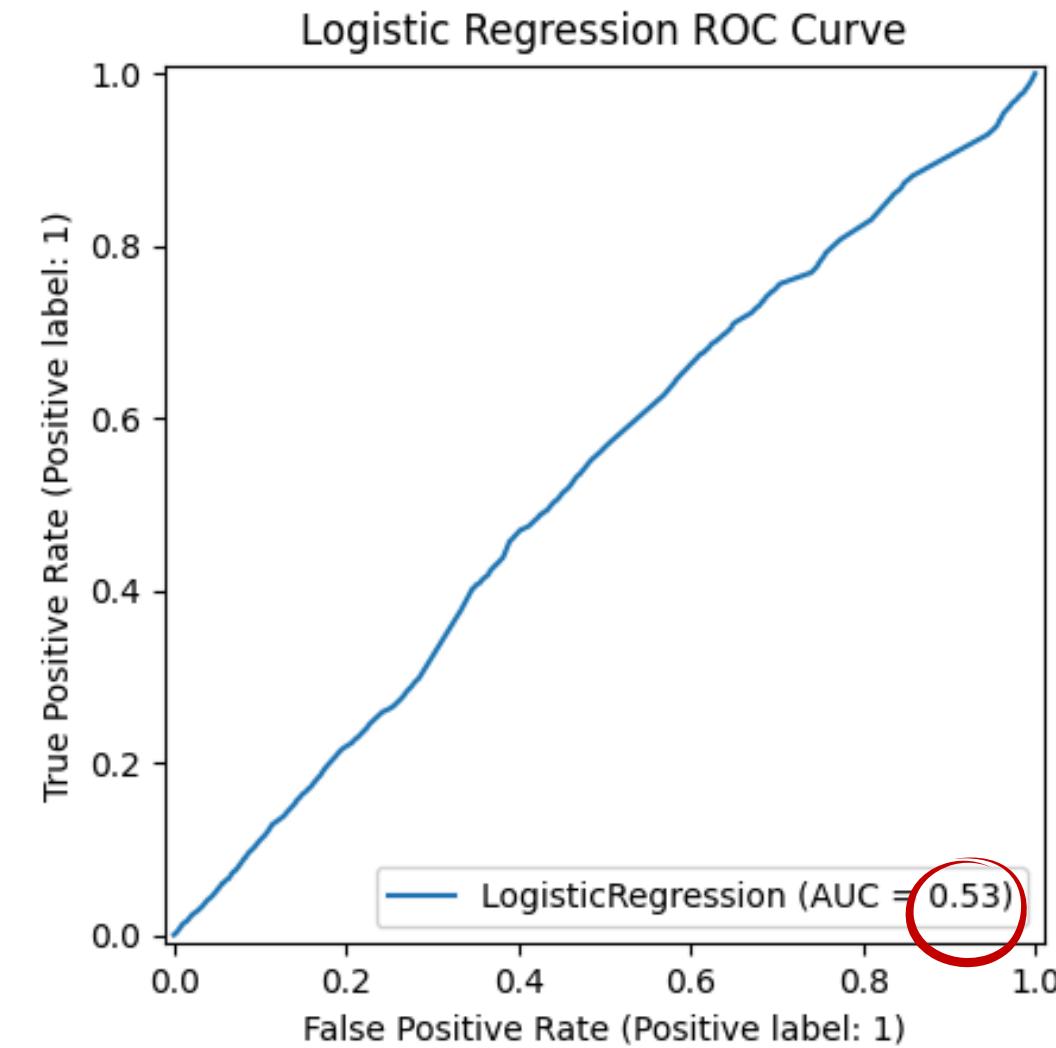
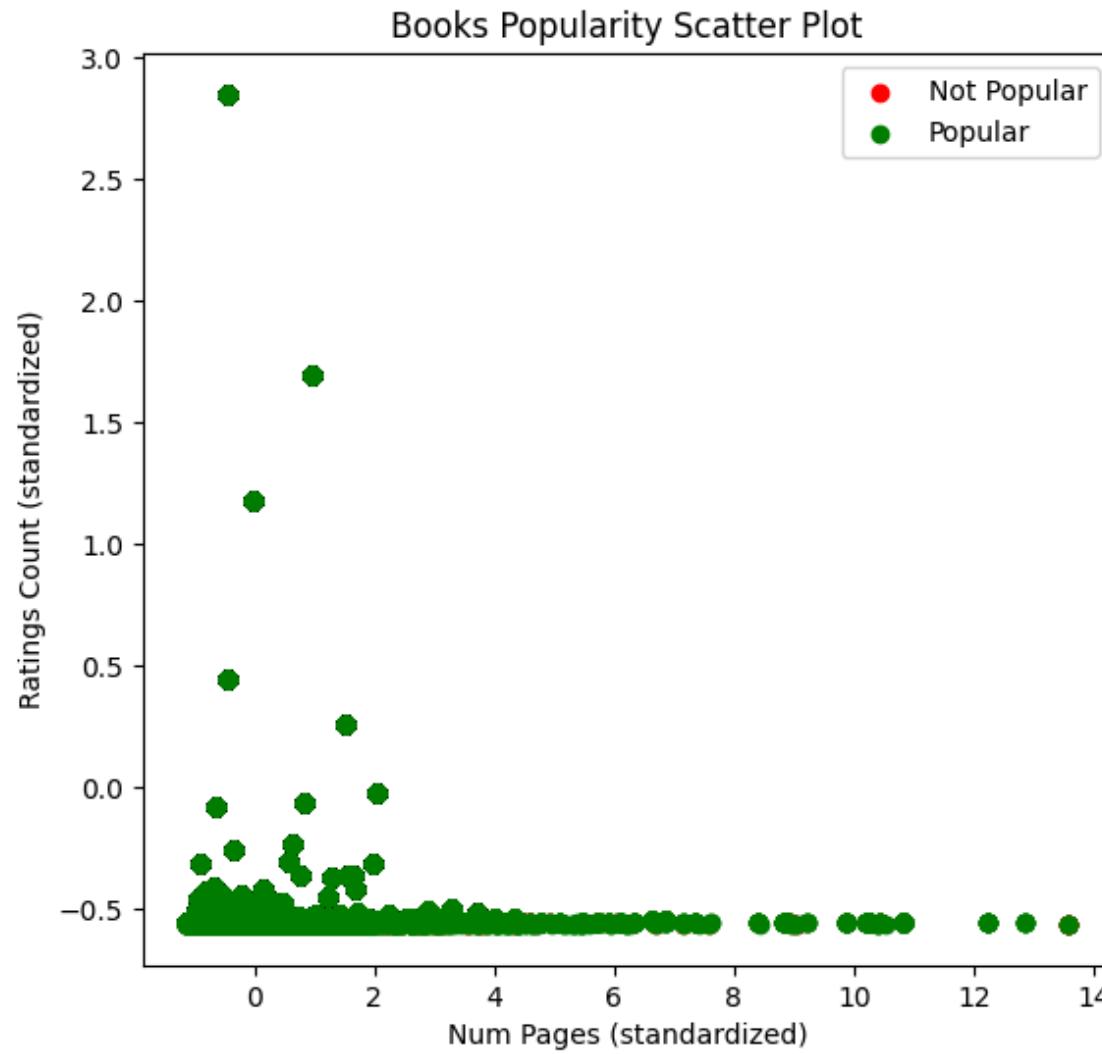
```
# Visualization: Scatter plot  
plt.figure(figsize=(8,6))  
plt.scatter(D_test[y_test==0][:,1], D_test[y_test==0][:,2], color='red', label='Not Popular')  
plt.scatter(D_test[y_test==1][:,1], D_test[y_test==1][:,2], color='green', label='Popular')  
plt.xlabel('Num Pages (standardized)')  
plt.ylabel('Ratings Count (standardized)')  
plt.legend()  
plt.title('Books Popularity Scatter Plot')  
plt.show()  
  
# ROC Curve  
RocCurveDisplay.from_estimator(model, D_test, y_test)  
plt.title("Logistic Regression ROC Curve")  
plt.show()
```

scatter plot

ROC curve

ROC curve* to evaluate classification performance.

try 1: Logistic Regression



Accuracy : 0.7548
Precision: 0.7548
Recall : 1.0000
F1-score : 0.8603
ROC-AUC : 0.5310

try 2: Random Forest

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (accuracy_score,precision_score,recall_score,f1_score,roc_auc_score,classification_report,RocCurveDisplay)

# Load dataset
dataset = pd.read_csv("data/processed/cleaned.csv") read the dataset

# Binary classification: books with rating > 3 are considered popular
popularityBound = 3
dataset["label"] = (dataset["rating"] > popularityBound).astype(int) classification

# Feature selection
features = ["publication_year","num_pages","ratings_count","average_rating"] feature selection

X = dataset[features]
y = dataset["label"]

# Train and Test split
X_train, X_test, y_train, y_test = train_test_split(X, y,test_size=0.3,random_state=42,stratify=y) dataset splitting
```

try 2: Random Forest

```
# Random Forest model
model = RandomForestClassifier(
    n_estimators=200,
    class_weight="balanced",
    random_state=42,
    n_jobs=-1
)

model.fit(X_train, y_train)

# Prediction and evaluation
y_pred = model.predict(X_test)
y_proba = model.predict_proba(X_test)[:, 1]

acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
auc = roc_auc_score(y_test, y_proba)

print("Random Forest Performance")
print("-" * 40)
print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall   : {rec:.4f}")
print(f"F1-score : {f1:.4f}")
print(f"ROC-AUC  : {auc:.4f})
```

random forest model

```
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

# Feature importance
importances = model.feature_importances_
feat_imp = pd.Series(importances, index=features).sort_values(ascending=False)

print("\nFeature Importances:")
print(feat_imp)

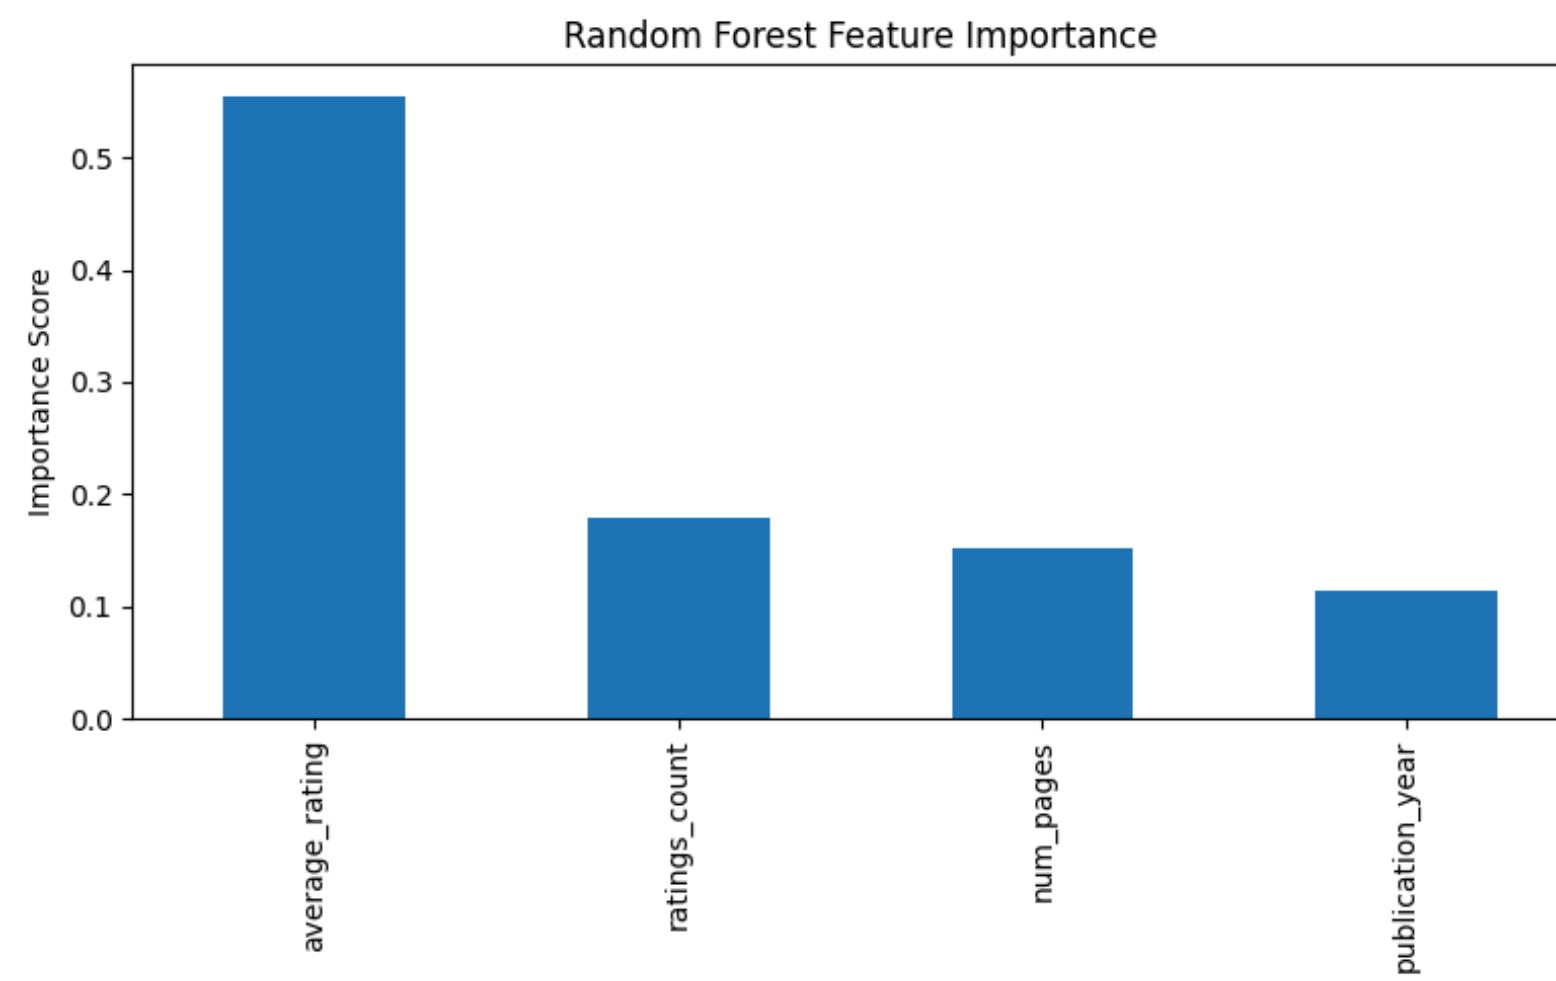
# Visualization
plt.figure(figsize=(8, 5))
feat_imp.plot(kind="bar")
plt.title("Random Forest Feature Importance")
plt.ylabel("Importance Score")
plt.tight_layout()
plt.show()

# Roc curve
RocCurveDisplay.from_estimator(model, X_test, y_test)
plt.title("Random Forest ROC Curve")
plt.show()
```

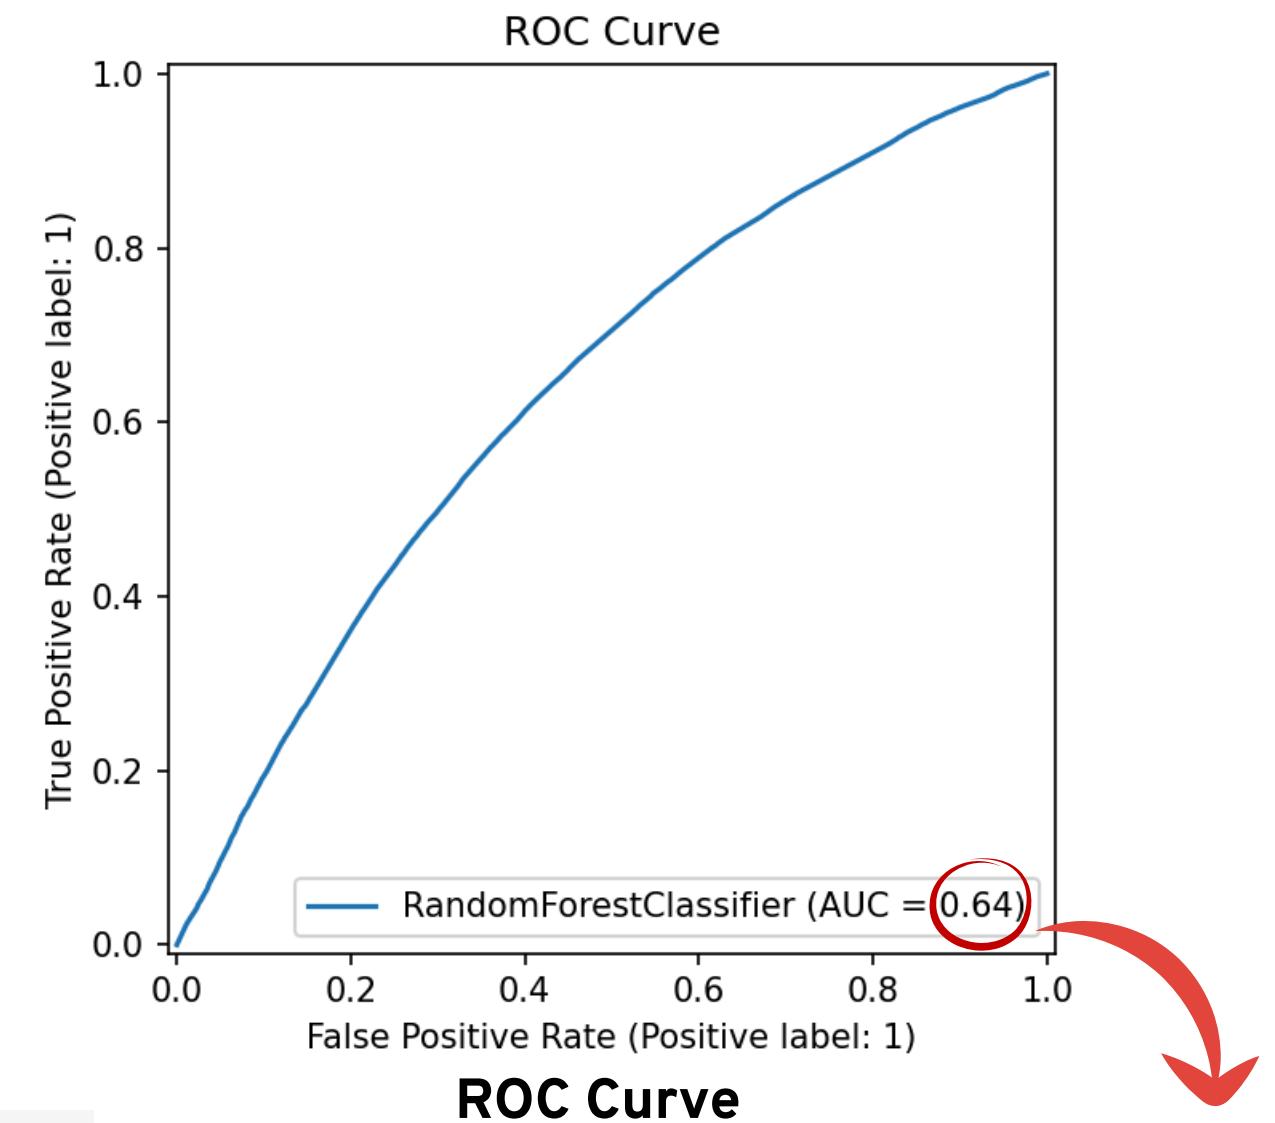
feature importance

prediction and accuracy

try 2: Random Forest



Accuracy : 0.6172
Precision: 0.8211
Recall : 0.6298
F1-score : 0.7128
ROC-AUC : 0.6376



Random Forest achieved a higher AUC value compared to Logistic Regression.

try 3 : Multilayer Perceptron

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.neural_network import MLPClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.model_selection import train_test_split
from imblearn.over_sampling import SMOTE
from sklearn.metrics import (accuracy_score,precision_score,recall_score,f1_score, roc_auc_score,classification_report,RocCurveDisplay)

# Read the dataset (relative path)
dataset = pd.read_csv("data/processed/cleaned.csv") read the dataset

# Define a rule for classification
popularityBound = 3
dataset['label'] = (dataset['rating'] > popularityBound).astype(int) classification

# Select features
features = ['publication_year', 'num_pages', 'ratings_count','average_rating']

D = dataset[features] # Feature matrix
y = dataset['label'] # Target variable feature selection

# Split the dataset into 70% train and 30% test
D_train, D_test, y_train, y_test = train_test_split( D, y,test_size=0.3,random_state=42,stratify=y) dataset splitting
```

try 3 : Multilayer Perceptron

```
# Scaler for feature normalization
scaler = StandardScaler()                                     normalization
D_train_scaled = scaler.fit_transform(D_train)
D_test_scaled = scaler.transform(D_test)

# Balance classes using SMOTE
smote = SMOTE(random_state=42, sampling_strategy=0.5)        SMOTE
D_train_res, y_train_res = smote.fit_resample(D_train_scaled, y_train)

# Define MLP as base learner
mlp_model = MLPClassifier(
    hidden_layer_sizes=(32,16,8), # Three hidden layers with decreasing size
    activation='relu',           # ReLU activation for hidden layers
    solver='adam',               # Adam optimizer
    learning_rate='adaptive',    # Adaptive learning rate
    max_iter=2000,               # Maximum training iterations
    early_stopping=True,         # Stop early if validation score does not improve
    random_state=42
)
```

SMOTE*: To improve the model, it generates synthetic samples for the minority class.
MLPClassifier* is a single neural network model that learns patterns from the data.

try 3 : Multilayer Perceptron

```
# Configure Bagging Ensemble  
bagging = BaggingClassifier(  
    estimator=mlp_model,  
    n_estimators=50,  
    max_samples=0.8,  
    max_features=1.0,  
    bootstrap=True,  
    n_jobs=-1,  
    random_state=42  
)  
  
# Train Bagging Ensemble  
bagging.fit(D_train_res, y_train_res)
```

ensemble model

training

We define **mlp_model** first as the **base learner** for the **bagging ensemble**.

BaggingClassifier creates multiple copies of **mlp_model** and trains each on a different subset of the data.

BaggingClassifier is an ensemble method that combines multiple base learners to reduce errors.

try 3 : Multilayer Perceptron

```
# Prediction and evaluation
y_pred = bagging.predict(D_test_scaled)
y_proba = bagging.predict_proba(D_test_scaled)[:, 1]

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
auc = roc_auc_score(y_test, y_proba)

print(f"Accuracy : {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall   : {recall:.4f}")
print(f"F1-score : {f1:.4f}")
print(f"ROC-AUC  : {auc:.4f}")

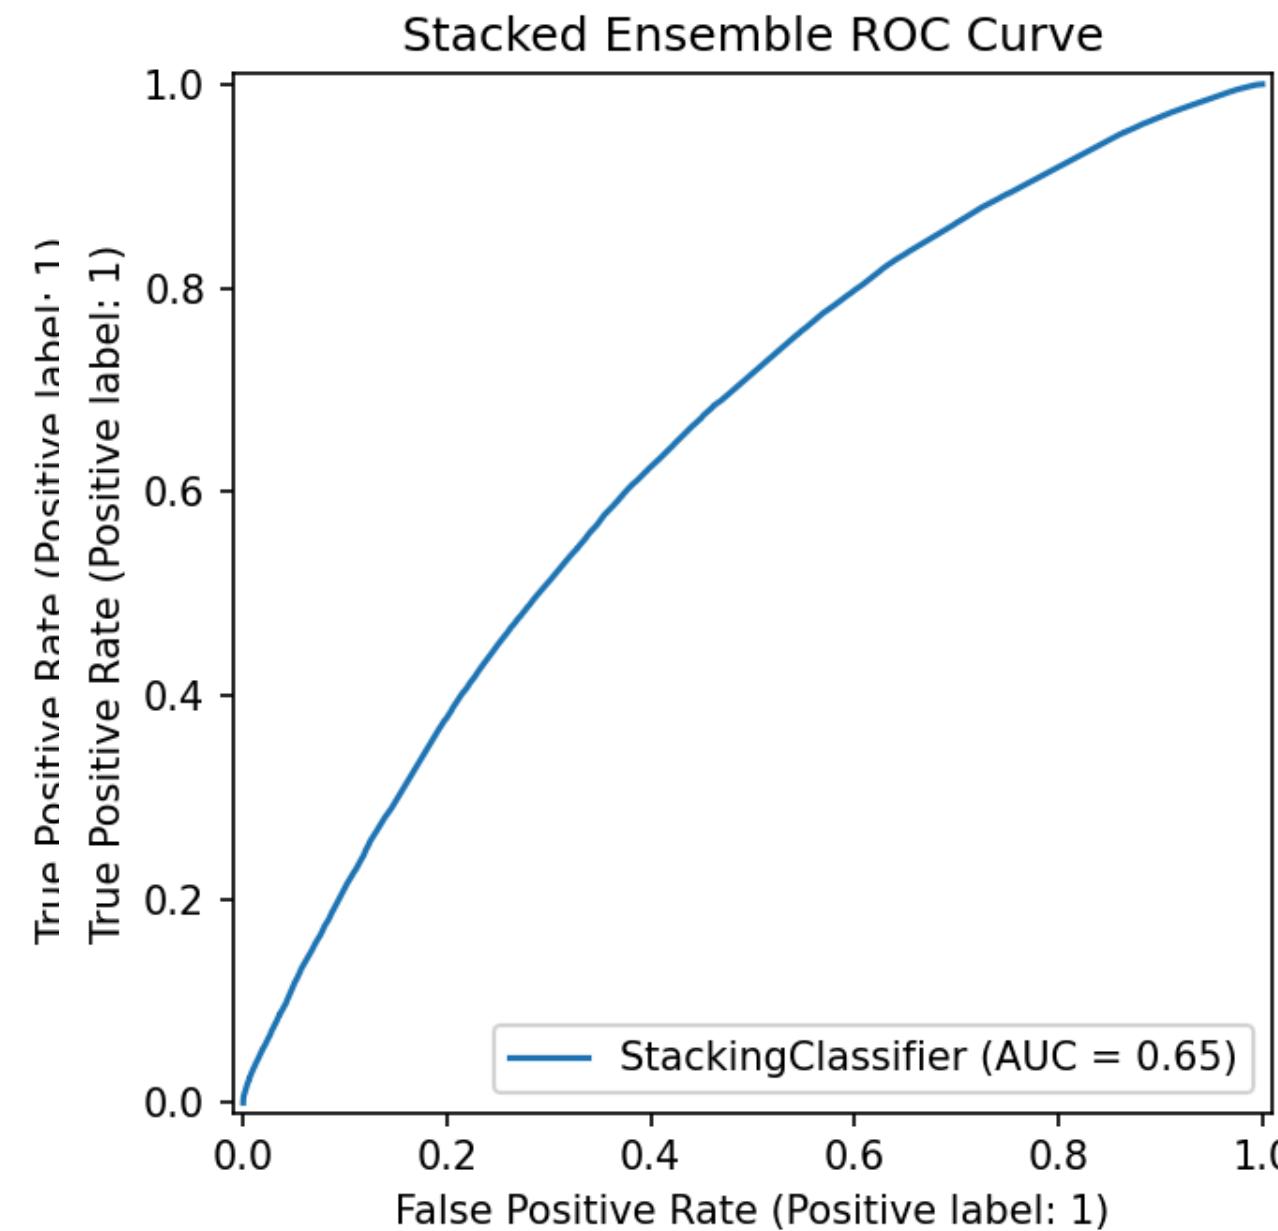
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

# ROC Curve
RocCurveDisplay.from_estimator(bagging, D_test_scaled, y_test)
plt.title("Optimized Bagging + MLP ROC Curve")
plt.show()
```

prediction and accuracy

ROC curve

try 3 : Multilayer Perceptron



Accuracy : 0.7506
Precision: 0.7741
Recall : 0.9453
F1-score : 0.8512
ROC-AUC : 0.6537

Our dataset has too few numeric features, so the MLP does not have enough information, making the model weak.

try 4 : Adaboost

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import (accuracy_score,precision_score,recall_score,f1_score,roc_auc_score,classification_report,RocCurveDisplay)

# Read the dataset (relative path)
dataset = pd.read_csv("data/processed/cleaned.csv") read the dataset

# Create binary popularity label
dataset['label'] = (dataset['rating'] > 3).astype(int) classification

# Select features
features = ['publication_year', 'num_pages', 'ratings_count', 'average_rating'] feature selection

D = dataset[features] # Feature matrix
y = dataset['label'] # Target variable

# Split the dataset into 70% train and 30% test
D_train, D_test, y_train, y_test = train_test_split(D, y,test_size=0.3,random_state=42,stratify=y) dataset splitting

# Scaler for feature normalization
scaler = StandardScaler()
D_train_scaled = scaler.fit_transform(D_train)
D_test_scaled = scaler.transform(D_test) normalization
```

try 4 : Adaboost

```
# Weak learner  
base_estimator = DecisionTreeClassifier(max_depth=1, random_state=42)  
  
# Configure AdaBoost with decision tree base learner  
adaboost_model = AdaBoostClassifier(  
    estimator=base_estimator,  
    n_estimators=100,  
    learning_rate=0.5, # Weight of each weak learner  
    random_state=42  
)  
  
# Train AdaBoost model  
adaboost_model.fit(D_train_scaled, y_train)
```

base model

AdaBoost model

training

Bagging trains **many independent models in parallel** using bootstrap samples, while **AdaBoost** trains models **sequentially** and increases the weight of **misclassified** samples to focus on hard cases.

base_estimator* : A shallow decision stump (single-split tree) used as a **fast weak learner for AdaBoost**. AdaBoost* **trains n weak learners and combines them**, giving more weight to hard examples.

try 4 : Adaboost

```
# Prediction and evaluation
y_pred = adaboost_model.predict(D_test_scaled)
y_proba = adaboost_model.predict_proba(D_test_scaled)[:, 1]

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
auc = roc_auc_score(y_test, y_proba)

print(f"Accuracy : {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall   : {recall:.4f}")
print(f"F1-score : {f1:.4f}")
print(f"ROC-AUC  : {auc:.4f}")

# classification report
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

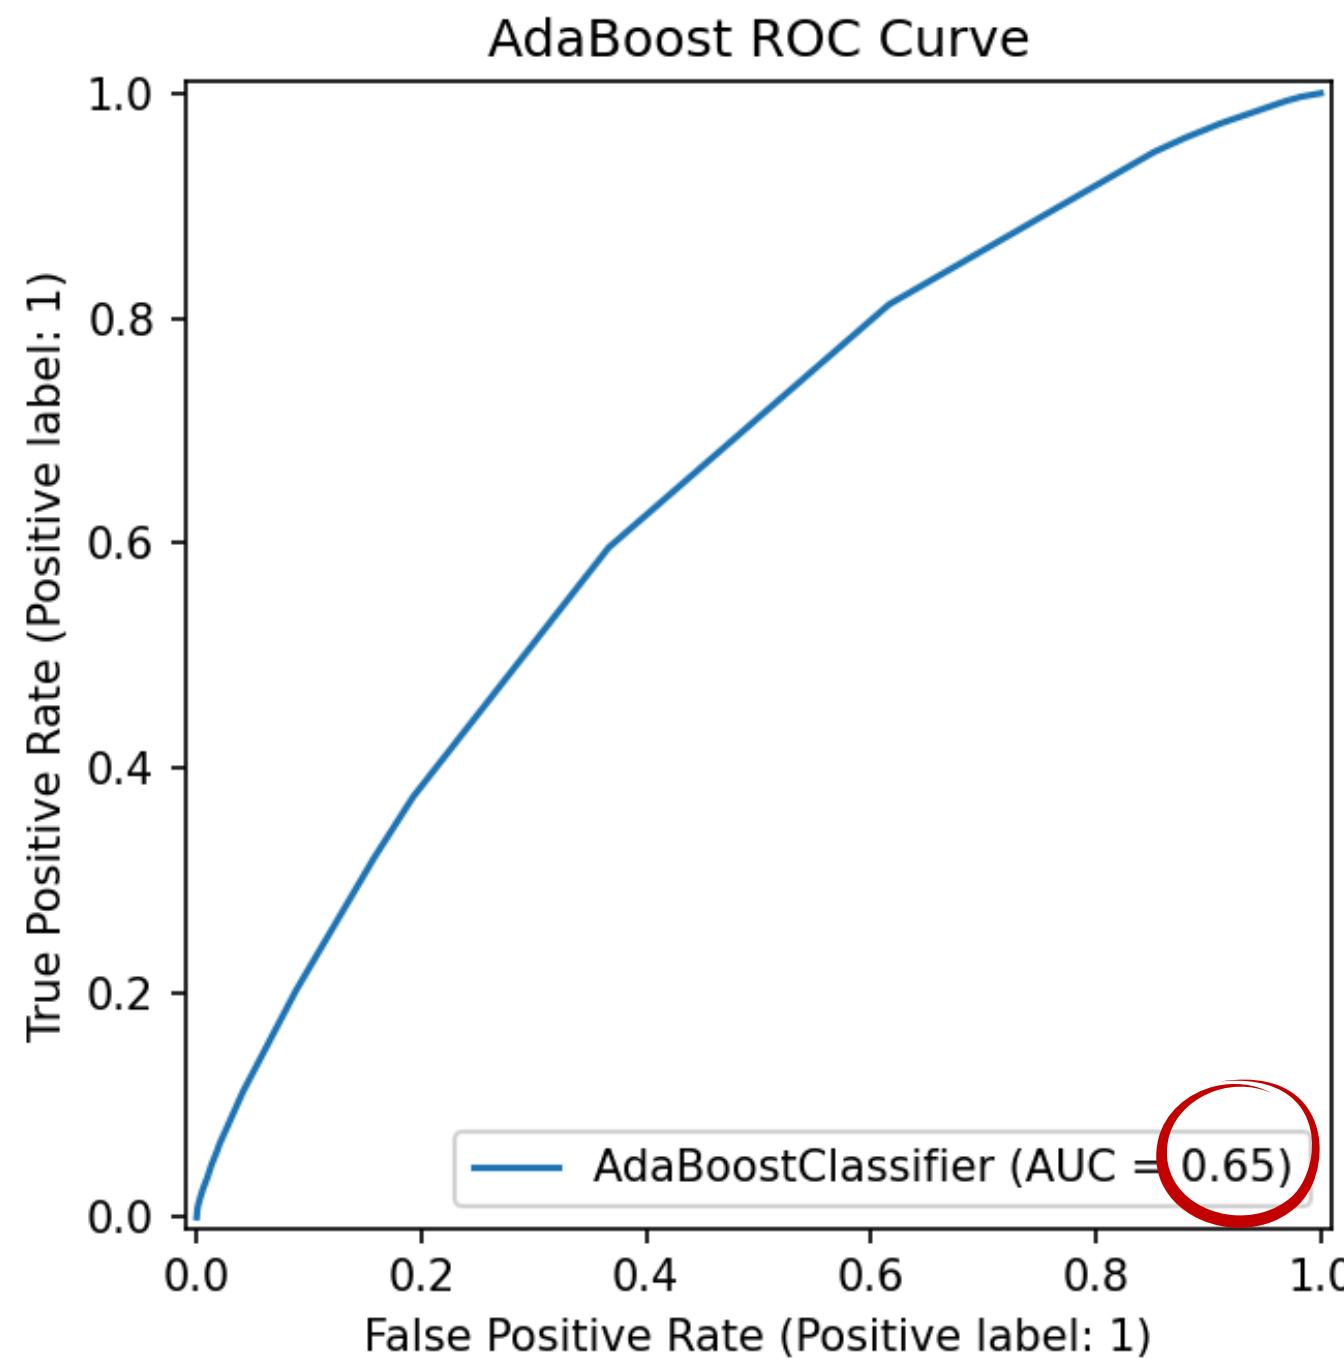
# ROC curve
RocCurveDisplay.from_estimator(adaboost_model, D_test_scaled, y_test)
plt.title("AdaBoost ROC Curve")
plt.show()
```

prediction and accuracy

ROC curve

try 4 : Adaboost

```
Accuracy : 0.7560
Precision: 0.7564
Recall   : 0.9979
F1-score : 0.8605
ROC-AUC  : 0.6496
```



try 5 : Stacking Ensemble Using Logistic Regression, Random Forest, and MLP

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, StackingClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.model_selection import train_test_split
from imblearn.over_sampling import SMOTE
from sklearn.metrics import (accuracy_score,precision_score,recall_score,f1_score,roc_auc_score,classification_report, RocCurveDisplay)

# Read the dataset (relative path)
dataset = pd.read_csv("data/processed/cleaned.csv")
```

read the dataset

```
# Define a rule for classification
popularityBound = 3
dataset['label'] = (dataset['rating'] > popularityBound).astype(int)
```

classification

```
# Select features
features = ['publication_year', 'num_pages', 'ratings_count', 'average_rating']

D = dataset[features] # Feature matrix
y = dataset['label'] # Target variable
```

feature selection

```
# Split dataset into 70% train and 30% test
D_train, D_test, y_train, y_test = train_test_split(D, y,test_size=0.3,random_state=42,stratify=y)
```

dataset splitting

try 5 : Stacking Ensemble Using Logistic Regression, Random Forest, and MLP

```
# Scaler for feature normalization
scaler = StandardScaler()
D_train_scaled = scaler.fit_transform(D_train)
D_test_scaled = scaler.transform(D_test)

# Balance classes using SMOTE
smote = SMOTE(random_state=42, sampling_strategy=0.5)
D_train_res, y_train_res = smote.fit_resample(D_train_scaled, y_train)

# Base models for Stacking
estimators = [
    ('lr', LogisticRegression(max_iter=1000, random_state=42)),
    ('rf', RandomForestClassifier(
        n_estimators=200,
        class_weight='balanced',
        random_state=42,
        n_jobs=-1
    )),
    ('mlp', MLPClassifier(
        hidden_layer_sizes=(32,16,8),
        activation='relu',
        solver='adam',
        learning_rate='adaptive',
        max_iter=2000,
        early_stopping=True,
        random_state=42
    ))
]

# Stacking Ensemble
stacked_model = StackingClassifier(
    estimators=estimators,
    final_estimator=LogisticRegression(max_iter=1000),
    n_jobs=-1
)

# Train model
stacked_model.fit(D_train_res, y_train_res)
```

SMOTE

base models

ensemble model

try 5 : Stacking Ensemble Using Logistic Regression, Random Forest, and MLP

```
# Prediction and evaluation
y_pred = stacked_model.predict(D_test_scaled)
y_proba = stacked_model.predict_proba(D_test_scaled)[:, 1]

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
auc = roc_auc_score(y_test, y_proba)

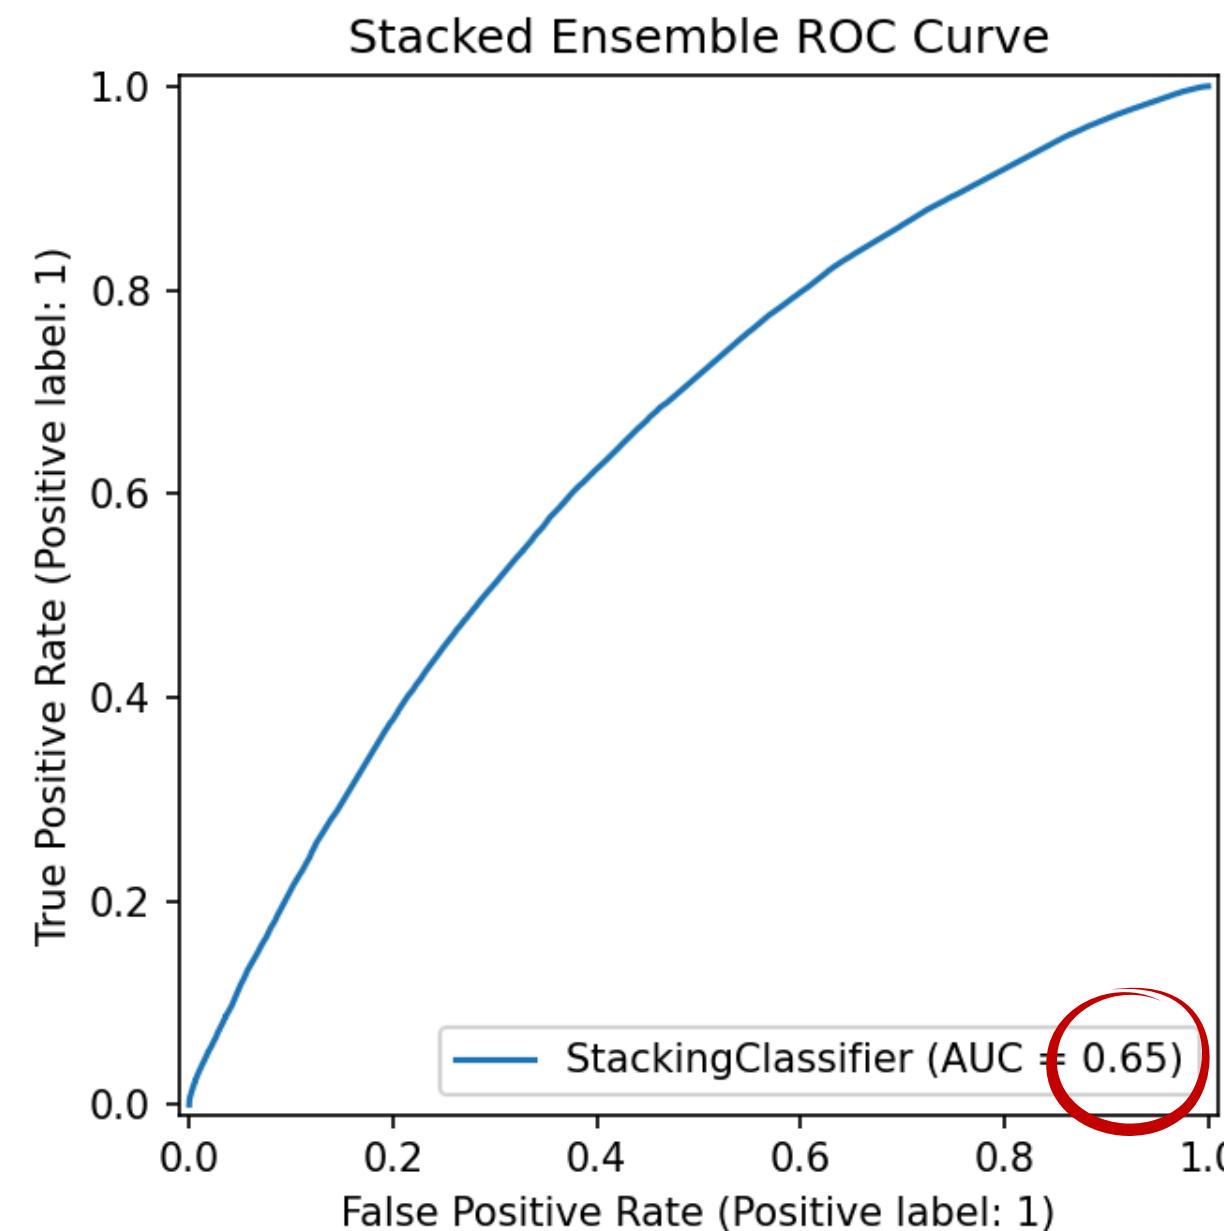
print(f"Accuracy : {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall   : {recall:.4f}")
print(f"F1-score : {f1:.4f}")
print(f"ROC-AUC  : {auc:.4f}")

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

# ROC Curve
RocCurveDisplay.from_estimator(stacked_model, D_test_scaled, y_test)
plt.title("Stacked Ensemble ROC Curve")
plt.show()
```

prediction and accuracy

try 5 : Stacking Ensemble Using Logistic Regression, Random Forest, and MLP



```
Accuracy : 0.7341
Precision: 0.7840
Recall   : 0.8937
F1-score : 0.8353
ROC-AUC  : 0.6493
```

try 6 : SVD

```
import pandas as pd
import numpy as np
from surprise import Dataset, Reader, SVD
from surprise.model_selection import train_test_split
from surprise import accuracy
import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, auc

# Read the dataset (relative path)
df = pd.read_csv("data/processed/cleaned.csv")

# Define binary popularity for ROC curve (rating > 3 as popular)
df['label'] = (df['rating'] > 3).astype(int)

# Create Surprise dataset
reader = Reader(rating_scale=(1,5)) # Define rating scale for Surprise
data = Dataset.load_from_df(df[['user_id', 'book_id', 'rating']], reader)

# Split the dataset into 70% train and 30% test
trainset, testset = train_test_split(data, test_size=0.3, random_state=42)

# Define and train SVD model
model = SVD(
    n_factors=50,          # Number of latent factors
    n_epochs=30,           # Number of training iterations
    lr_all=0.005,          # Learning rate for all parameters
    reg_all=0.02,           # Regularization term
    random_state=42
)
```

read the dataset

surprise dataset

dataset splitting

SVD model

The SVD model* is a **recommendation system algorithm** that learns the hidden (latent) features of users and items to predict missing ratings.

try 6 : SVD

```
model.fit(trainset)

# Prediction and evaluation
predictions = model.test(testset)
rmse = accuracy.rmse(predictions) # Root Mean Squared Error
mae = accuracy.mae(predictions) # Mean Absolute Error
print(f" SVD RMSE: {rmse}, MAE: {mae}")

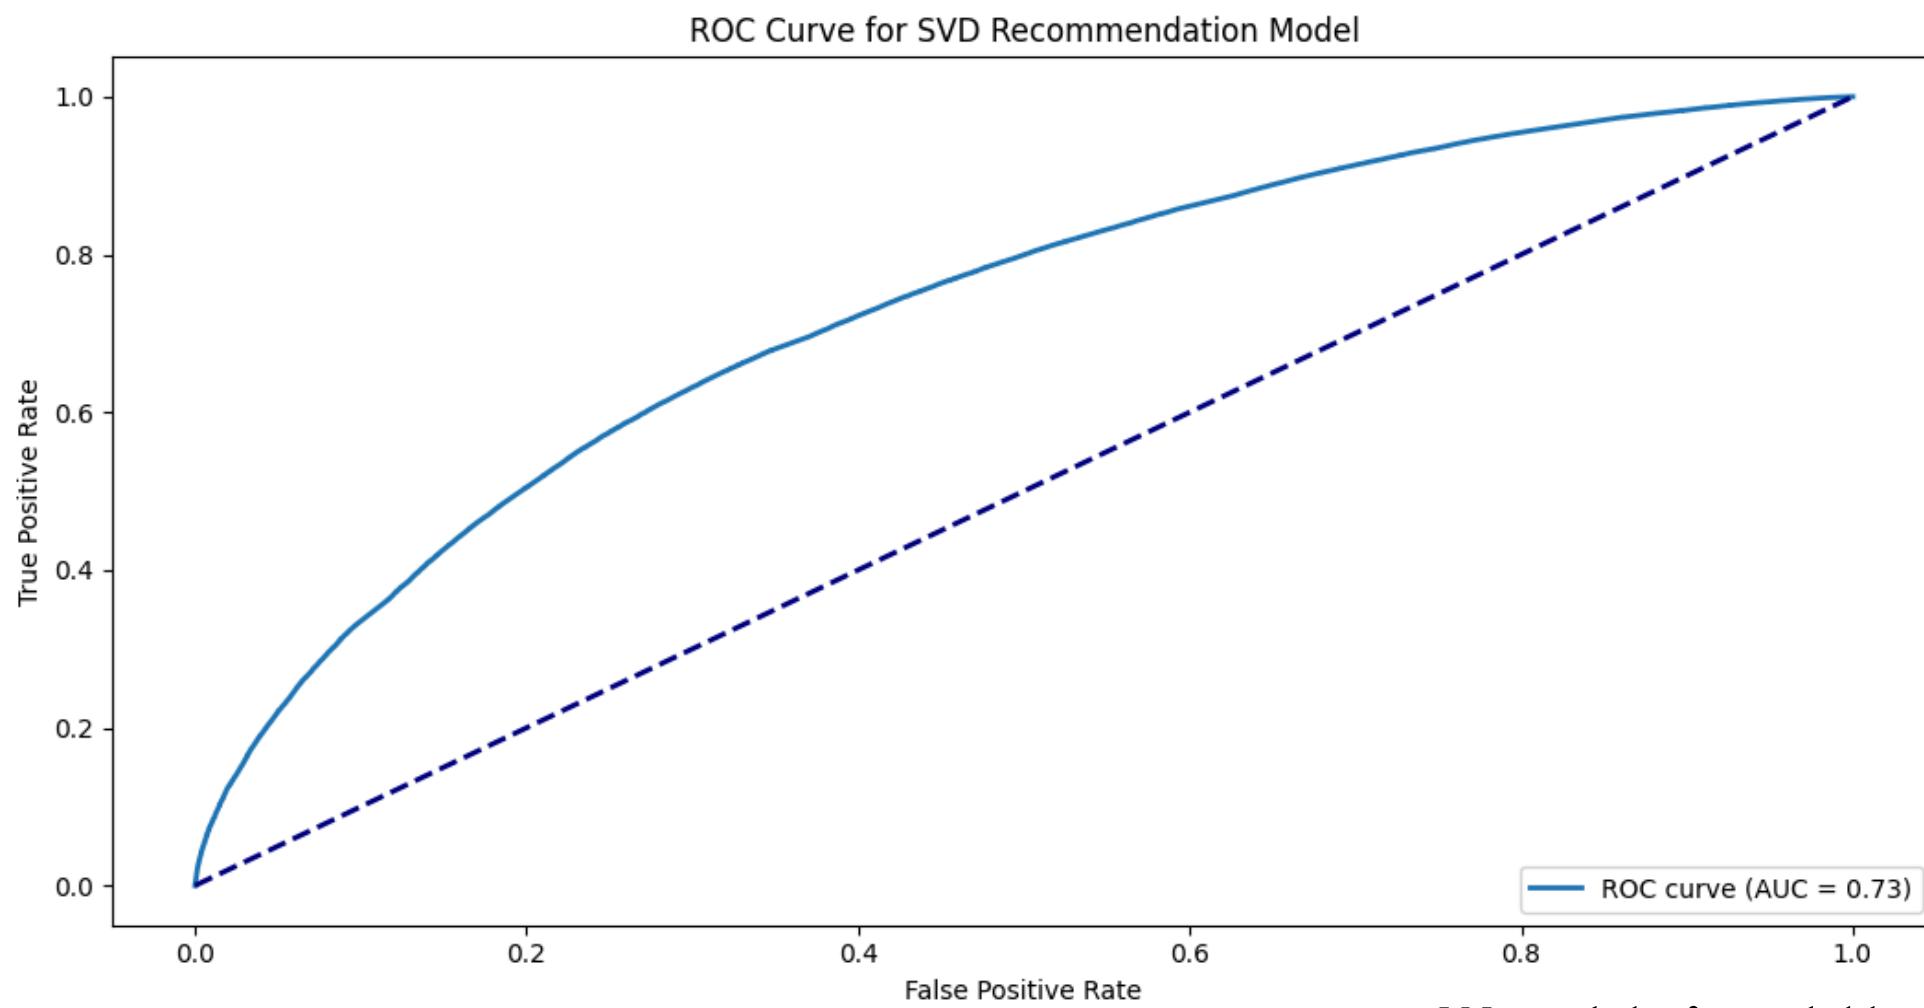
# ROC Curve preparation
y_true = np.array([pred.r_ui > 3 for pred in predictions]) # Binary labels
y_score = np.array([pred.est for pred in predictions]) # Predicted ratings
fpr, tpr, thresholds = roc_curve(y_true, y_score)
roc_auc = auc(fpr, tpr)
print(f"SVD ROC-AUC: {roc_auc:.4f}")

# ROC Curve
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, lw=2, label=f'ROC curve (AUC = {roc_auc:.2f})')
plt.plot([0,1], [0,1], color='navy', lw=2, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for SVD Recommendation Model')
plt.legend(loc="lower right")
plt.show()
```

prediction and accuracy

ROC curve

try 6 : SVD



We obtained the best AUC result with the SVD model, so we will proceed using SVD for the recommendation system.

RMSE: 0.8657

MAE: 0.6842

SVD RMSE: 0.8656942182728519, MAE: 0.6842251631518097

SVD ROC-AUC: 0.7262

Book
Recommendation
System with SVD

Book Recommendation System with SVD

```
import pandas as pd
import random
import pickle
from surprise import SVD

class SVDRecommender:
    """
    This class loads a trained SVD recommendation model
    and generates personalized book recommendations for users.
    """

    def __init__(self, model_path, data_path):
        # load the trained model
        with open(model_path, "rb") as f:
            self.model = pickle.load(f)

        # load data
        self.df = pd.read_csv(data_path)

        # cache for speed
        self.all_books = self.df['book_id'].unique() # all book's id in dataset
        self.all_users = self.df['user_id'].unique()

    # choice a random user function
    def get_random_user(self):
        return random.choice(self.all_users)
```

Book Recommendation System with SVD

```
# recommendation function
def recommend(self, user_id, top_n=10):
    # handle cold-start users
    if user_id not in self.all_users:
        raise ValueError("User ID not found in dataset.")

    seen_books = self.df[self.df['user_id'] == user_id]['book_id'].tolist() # holds user's book
    candidate_books = [b for b in self.all_books if b not in seen_books]

    predictions = [] # this list holds predictions for each 'not seen' book
    for book in candidate_books:
        pred = self.model.predict(user_id, book)
        predictions.append((pred.iid, pred.est))

    # sort top_n books in descending order
    top_predictions = sorted(
        predictions,
        key=lambda x: x[1],
        reverse=True
    )[:top_n]

    # create a dataframe for HTML
    rec_df = pd.DataFrame(
        top_predictions,
        columns=['book_id', 'predicted_rating']
    )

    # format ratings for better display
    rec_df['predicted_rating'] = rec_df['predicted_rating'].round(2)

    return rec_df[['title', 'predicted_rating']]
```

seen_books holds the books that have already been rated by the user (user_id).

All books in the DataFrame are stored in all_books.

candidate_books contains all books that the user has not rated yet.

predictions contains the estimated ratings for each book in candidate_books predicted by the SVD model.

top_predictions contains the highest predicted ratings from predictions, sorted in descending order, limited to top_n books.

recommended extracts the book IDs and their predicted ratings from top_predictions to prepare for creating a DataFrame with titles.

Book Recommendation System with SVD

Book Recommendation System

Personalized recommendations using SVD

Random User ID: 9a74aabe319e4c47dfd613d2d5c5b877

Book Title	Predicted Rating
Playing With Fire	4.94
A Great Big Ugly Man Came Up and Tied His Horse to Me: A Book of Nonsense Verse	4.93
Counting Descent	4.93
The Divan of Hafez in Original Persian	4.91
Maha Prasthanam	4.89

[Recommend Another User](#)

The model predicts that the user is highly likely to enjoy these books, as indicated by the consistently strong estimated ratings.